





DATA TOBOGGAN

What's New and Cool in Azure Synapse Analytics

Pawel Potasinski
Senior Program Manager

 @pawelpotasinski
 /in/pawelpotasinski




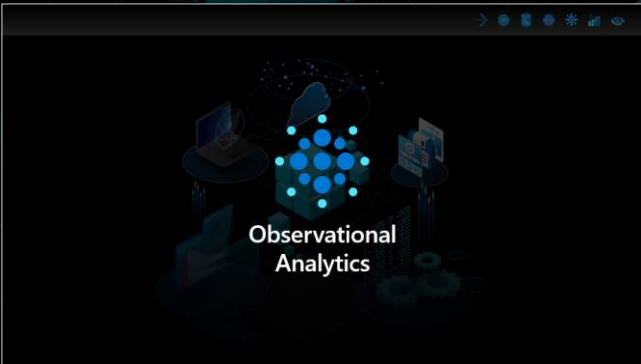
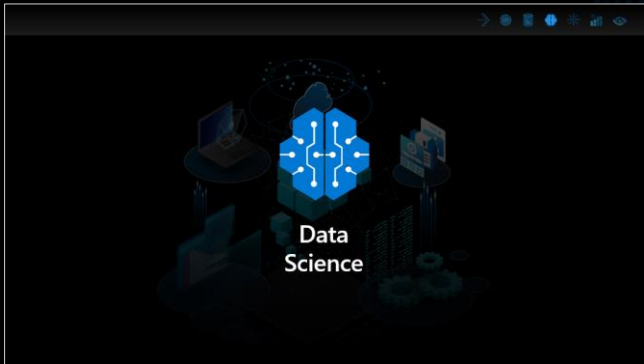


Stay up to date with Azure Synapse Analytics:

aka.ms/SynapseMonthlyUpdate

Synapse + Power BI







Data Integration

Data Integration

April-June 2022

- **[Azure Synapse Link for SQL \[Public Preview\]](#)**
- [Dataverse connector added to Synapse Data Flows](#)
- [Synapse Pipelines Web activity response timeout improvement](#)
- [Export pipeline monitoring as a CSV](#)
- [Incremental data loading made easy for Synapse and Azure Database for PostgreSQL and MySQL](#)
- [Assert Error Handling](#)
- [Mapping data flows projection editing](#)
- **[Fuzzy Join option in Join Transformation](#)**
- [Map Data \[Generally Available\]](#)
- **[Rerun pipeline with new parameters](#)**
- [User Defined Functions \[Generally Available\]](#)

Public Preview

Synapse Link for Microsoft SQL

Near real-time operational
analytics in Synapse

No data pipelines required


Hybrid integration for SQL
Server running on-premise or
in other clouds



SQL Server 2022
Public Preview Q2 2022



Azure SQL Database
Public Preview Q2 2022



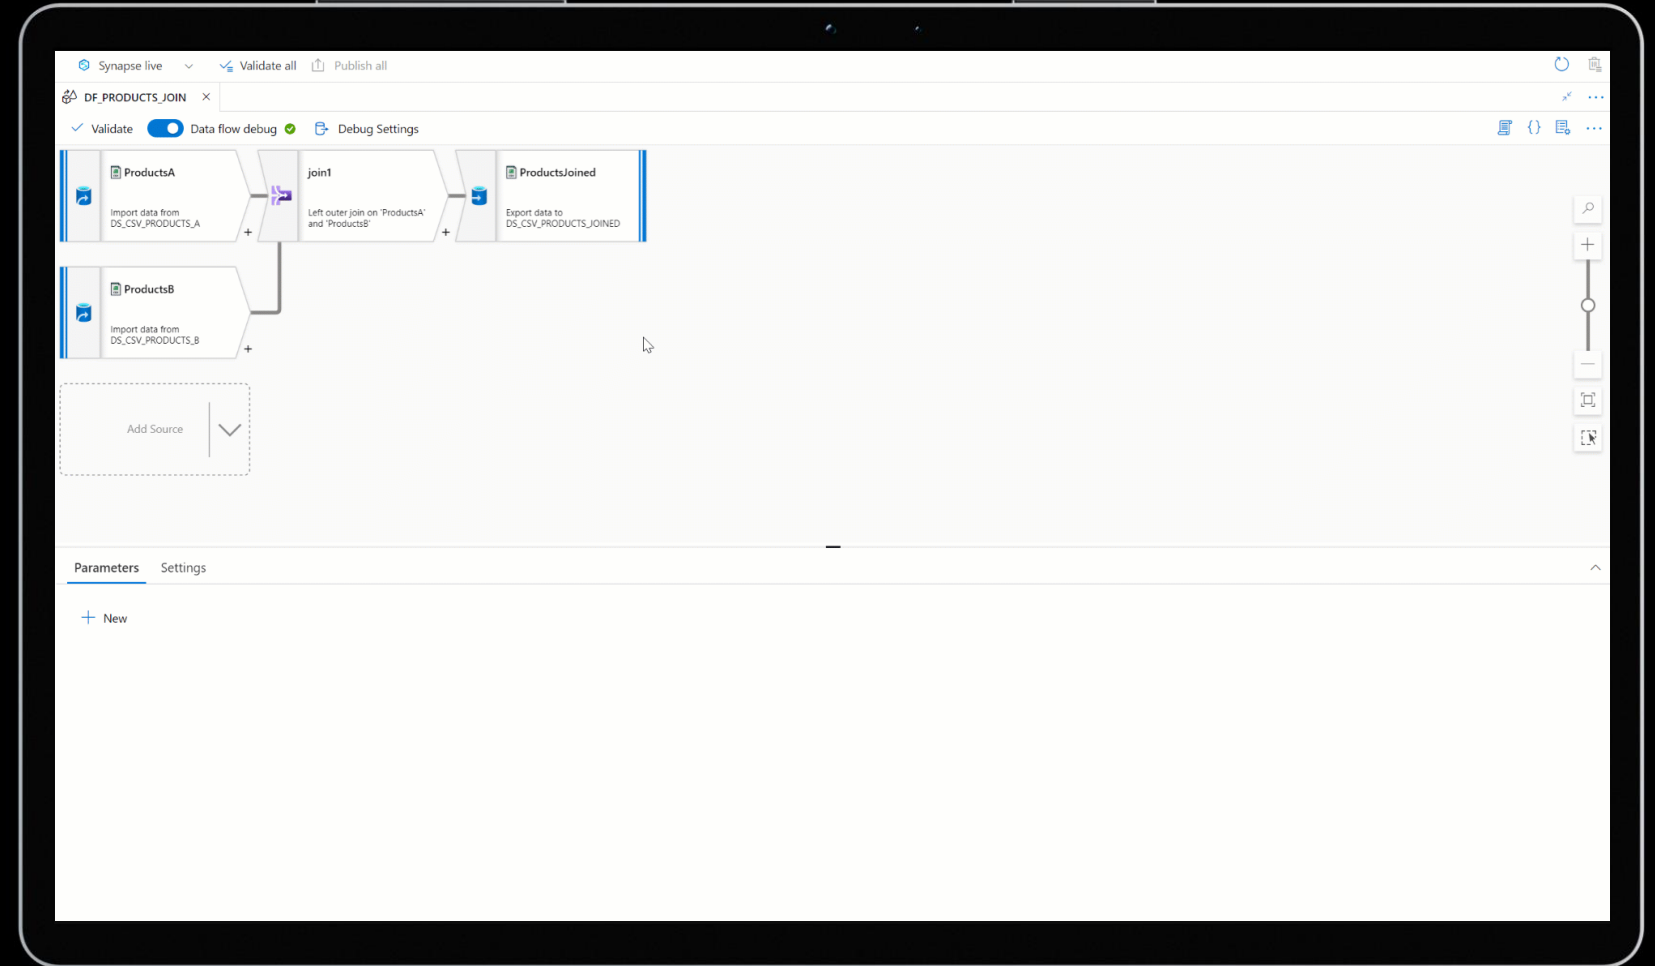
Fuzzy Join

No-code fuzzy logic for joins

Configurable similarity
threshold

Similarity score column added

Combine text parts



Rerun pipeline with new parameters

You can change parameter values when re-running a pipeline from the Monitoring page

Considered a new pipeline run

Pipeline runs

Triggered

Debug

Rerun

Cancel

Refresh

Edit columns

List

Gantt

Filter by run ID or name

Local time : Last 30 days

Pipeline name : All

Status : All

Runs : Latest runs

Triggered by : All

Add filter

Copy filters

Export to CSV

1 selected

Open in new tab

Showing 1 - 23 items

Pipeline name

pipeline2

pipeline2

pipeline2

pipeline2

pipeline2

pipeline2

pipeline2

pipeline2

pipeline2

pipeline2

Run start

Run end

Duration

Triggered by

Status

Error

00:00:02

00:00:02

00:00:02

00:00:02

00:00:01

00:00:02

00:00:02

00:00:02

00:00:01

00:00:02

00:00:02

JoshScheduled

JoshScheduled

JoshScheduled

JoshScheduled

JoshScheduled

JoshScheduled

JoshScheduled

JoshScheduled

JoshScheduled

JoshScheduled

JoshScheduled

Succeeded

Succeeded

Succeeded

Succeeded

Succeeded

Succeeded

Succeeded

Succeeded

Succeeded

Succeeded

Succeeded

Rerun?

Are you sure you want to rerun pipeline pipeline2?

OK

New parameters

Cancel

Pre-Order


New Book Alert!

Build scalable ETL jobs

Apply best practices for complex data pipelines

Cloud-based ETL patterns

Repeatable ETL design patterns



[Mapping Data Flows in Azure Data Factory: Building Scalable ETL Projects in the Microsoft Cloud by Mark Kromer at Amazon.com](#)



Data Engineering

Data Engineering

April-June 2022

- [**Apache Spark 3.2 \[Public Preview\]**](#)
- [Parameterization for Spark job definition](#)
- [Notebook snapshot](#)
- [Azure Synapse Dedicated SQL Pool Connector for Apache Spark Now Available in Python](#)
- [**Manage Azure Synapse Apache Spark configuration**](#)
- [**Reference unpublished notebooks**](#)

Public Preview

Spark 3.2

Enables developers
can leverage the latest
innovations in the
Spark ecosystem

Pandas (Koalas) integration

A highly popular and flexible library with broad industry adoption

Adaptive Query Execution (AQE) enabled by default

Significant improvements in query performance out-of-the-box

Small Query execution improvements

Small queries run faster due to reduced initialization overhead

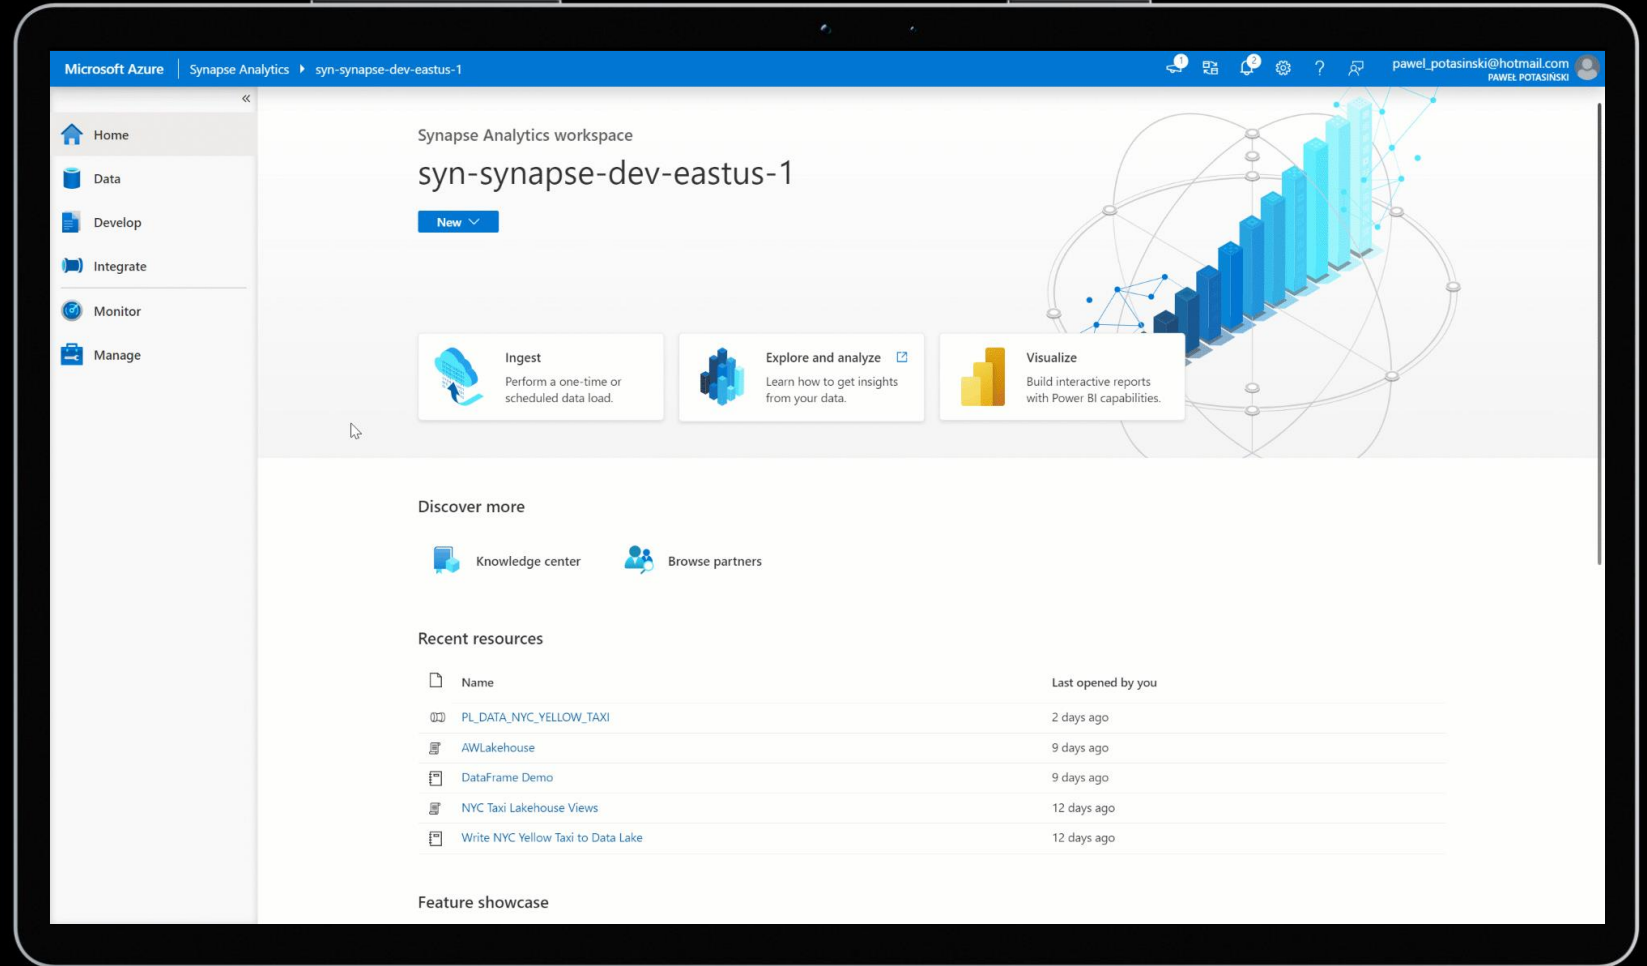
UPCOMING EVENT! [Register today](#) for our **Ask the Experts (ATE)** session with **Simon Whiteley on Spark 3.2** in Azure Synapse scheduled for **July 13 (10AM PDT)**. Learn more: <https://aka.ms/SynapseInfluencers>

Spark Configuration

Spark configuration artifact

Auto-suggestions and built-in
validation rules

Can be associated with a
Spark pool or a session,
re-used and exported

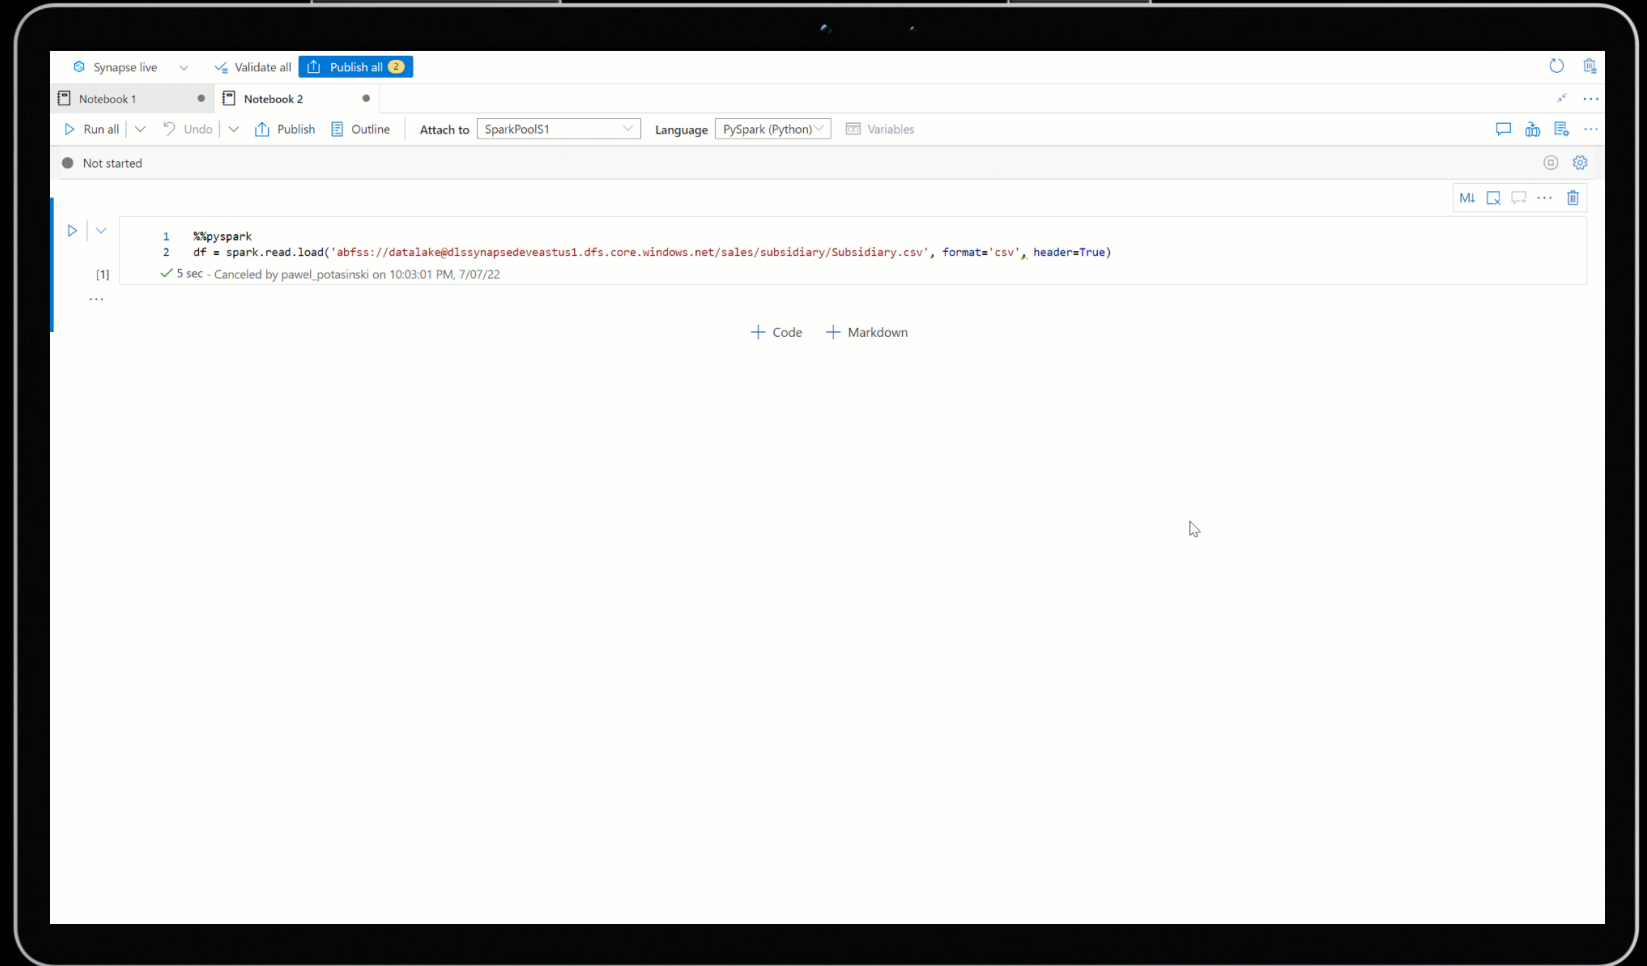


Reference unpublished notebooks

%run magic command can
work with unpublished
notebook

References notebook content
in web cache

Great for local debugging





Data Warehousing

Data Warehousing

April-June 2022

- **Cross-subscription restore for Azure Synapse SQL [Generally Available]**
- Recover SQL pool from dropped server or workspace
- Revamped exploration experience
- **Clone lake database**
- **Use wildcards to specify custom folder hierarchies**
- Data Warehouse Migration guide for Dedicated SQL Pools in Azure Synapse Analytics
- **Specify character column lengths... Not anymore!**
- Result set size limit increase


Public Preview

Cross subscription restore

Restore data warehouse across subscription boundary in one command without workarounds

Unlock Dev/Test scenarios and simplify billing at the subscription level

PowerShell and API support available



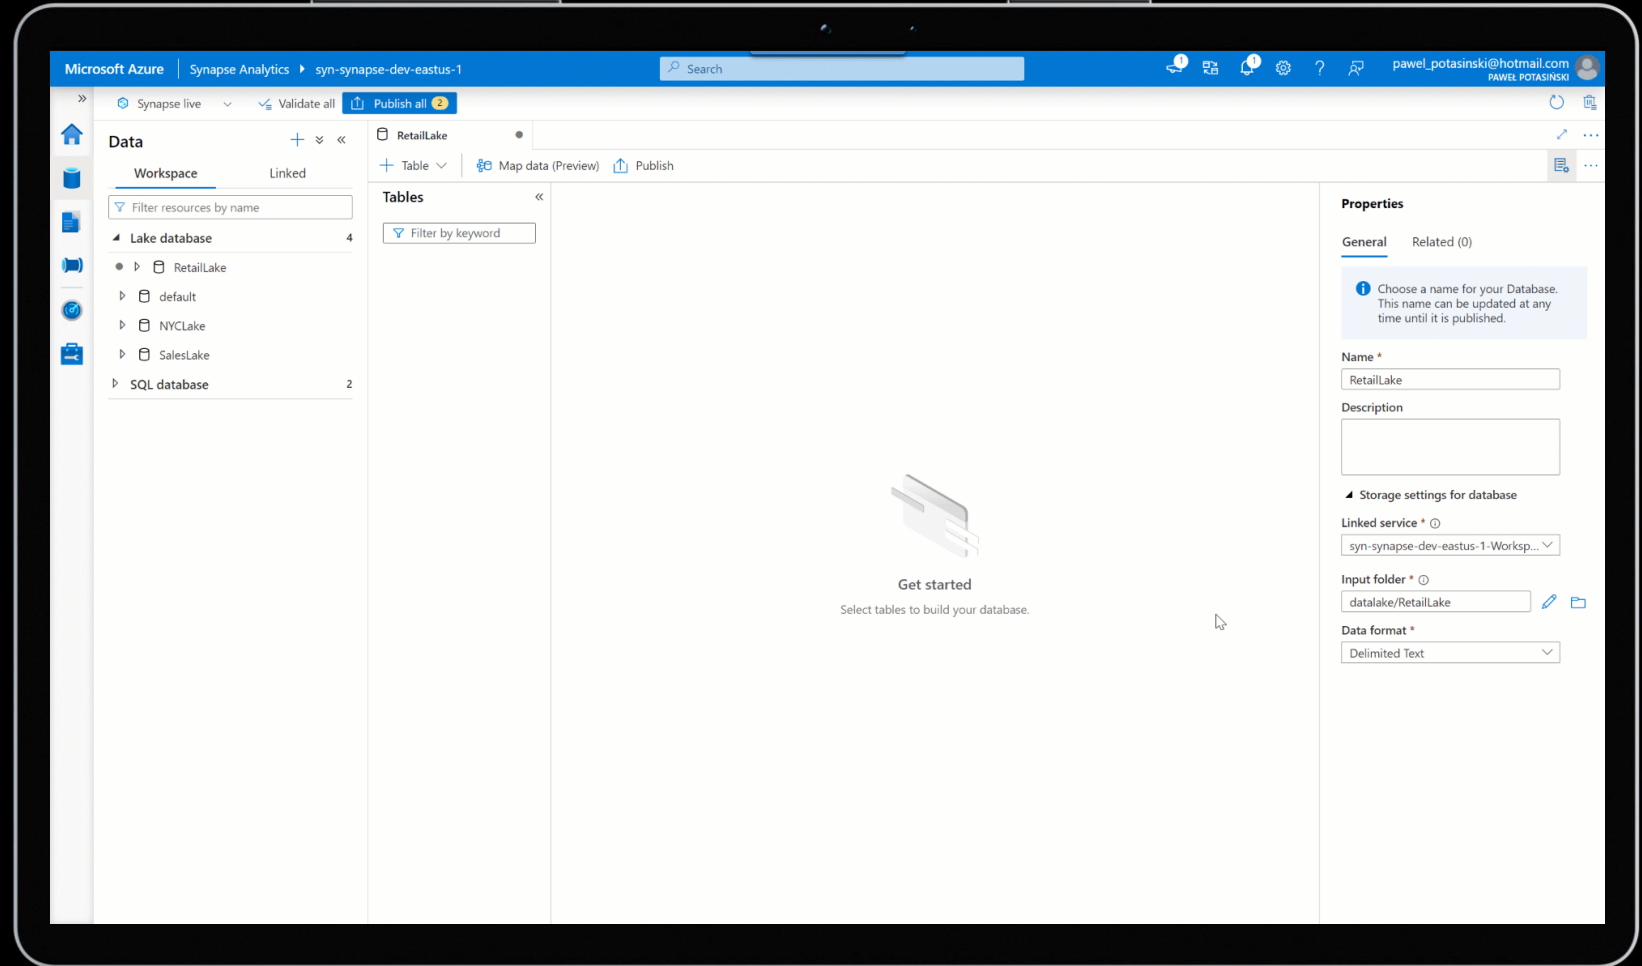
Generally Available

Database Templates & Database Designer

Built-in database templates

Low code database designer

Wildcards for custom folder
hierarchies

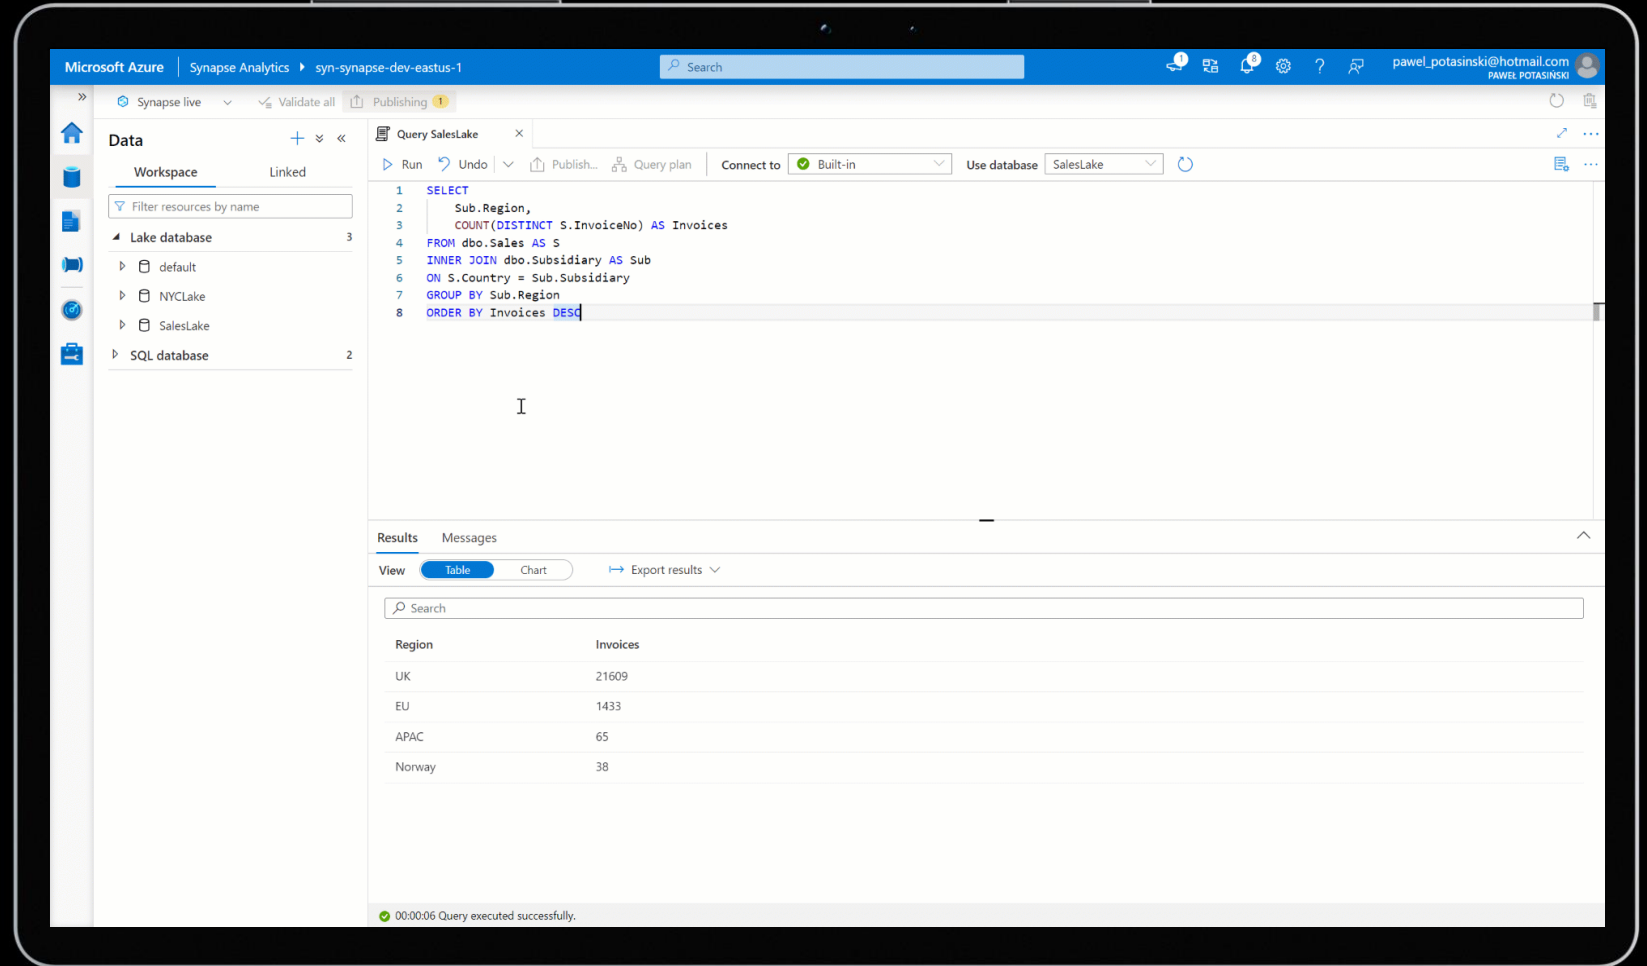


Clone lake database

Quick cloning of lake
databases and tables using
database designer UI

Editable metadata

No-code deployment




Character column lengths not required

No explicit schema definition required for performance

Works for OPENROWSET and external tables

Triggers whenever you omit schema or specify character column longer than 100 bytes

```
SELECT *
FROM
OPENROWSET (
    BULK 'https://pandemicdatalake.blob.core.windows.net/public/curated/covid-19/ecdc_cases/latest/ecdc_cases.csv',
    FORMAT = 'CSV',
    PARSER_VERSION = '2.0',
    HEADER_ROW = TRUE
) AS [r]
WITH (
    countries_and_territories VARCHAR(50)
) AS [r]
```



Data Science

April-June 2022

- [Distributed Deep Neural Network Training with Horovod and Petastorm \[Public Preview\]](#)

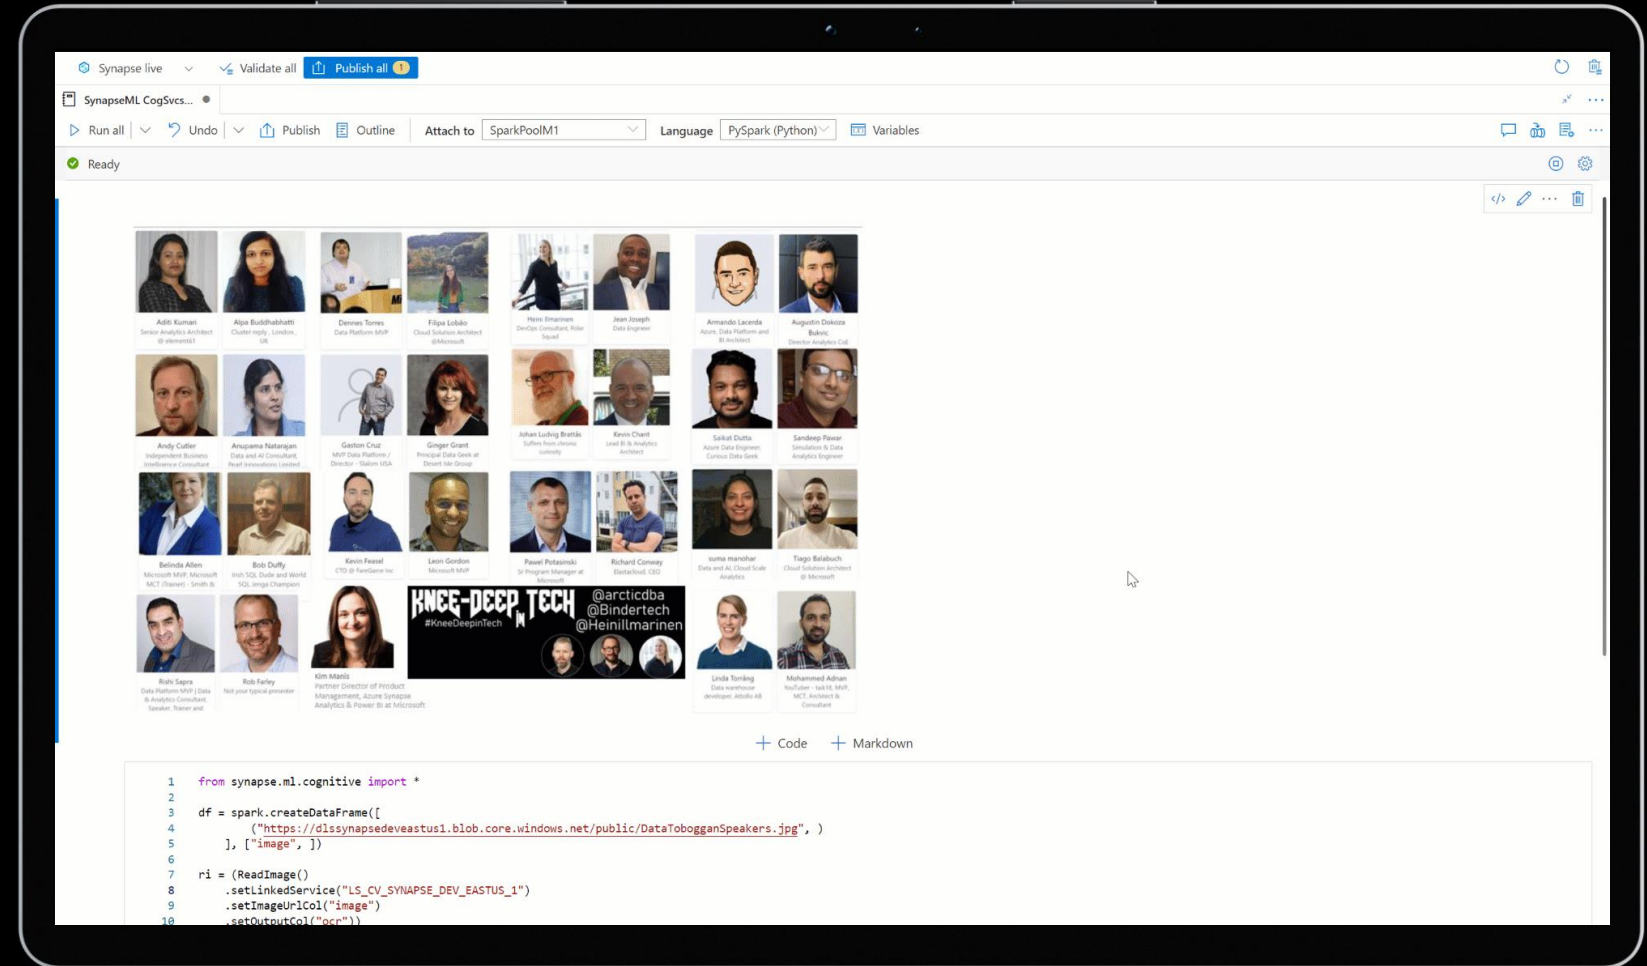
Generally Available

SynapseML


Open-Source Spark Library
Distributed ML algorithms for
Apache Spark

Cognitive Services Integration
for Spark

Language Binding Generators:
Python, R, Java (, .NET)



<https://aka.ms/spark>



Observational Analytics

Observational Analytics

April-June 2022

- [Synapse Data Explorer live query in Excel](#)
- [Use Managed Identities for External SQL Server Tables](#)
- [New KQL Learn module \(2 out of 3\) is live!](#)
- [Azure Synapse Data Explorer connector for Microsoft Power Automate, Logic Apps, and Power Apps](#)
- [Dynamic events routing from event hub to multiple databases](#)
- [Configure a database using a KQL inline script as part of JSON ARM deployment template](#)
- [Web Explorer new homepage](#)
- [Web Explorer sample gallery](#)
- [Web Explorer dashboards drill through capabilities](#)
- [Time Zone settings for Web Explorer](#)



Live query in Excel

Web explorer can create an
Excel workbook with
connection to the cluster

KQL executed by PowerQuery

Azure Data Explorer

Home

Data

Query

Dashboards (Preview)

My cluster (Preview)

MyFreeCluster

Test

StormEvents

help

ContosoSales

SampleTData

SampleLogs

SampleMetrics

Samples

MyFreeCluster.Test

help.SampleTData

help.SampleTData

MyFreeCluster.Test

Filter...

Run

Recall

Scope: @MyFreeCluster/Test

1 StormEvents

2 | sort by StartTime desc

3 | take 10

Table 1

Stats

Search

UTC

Done (0.767 s)

10 records

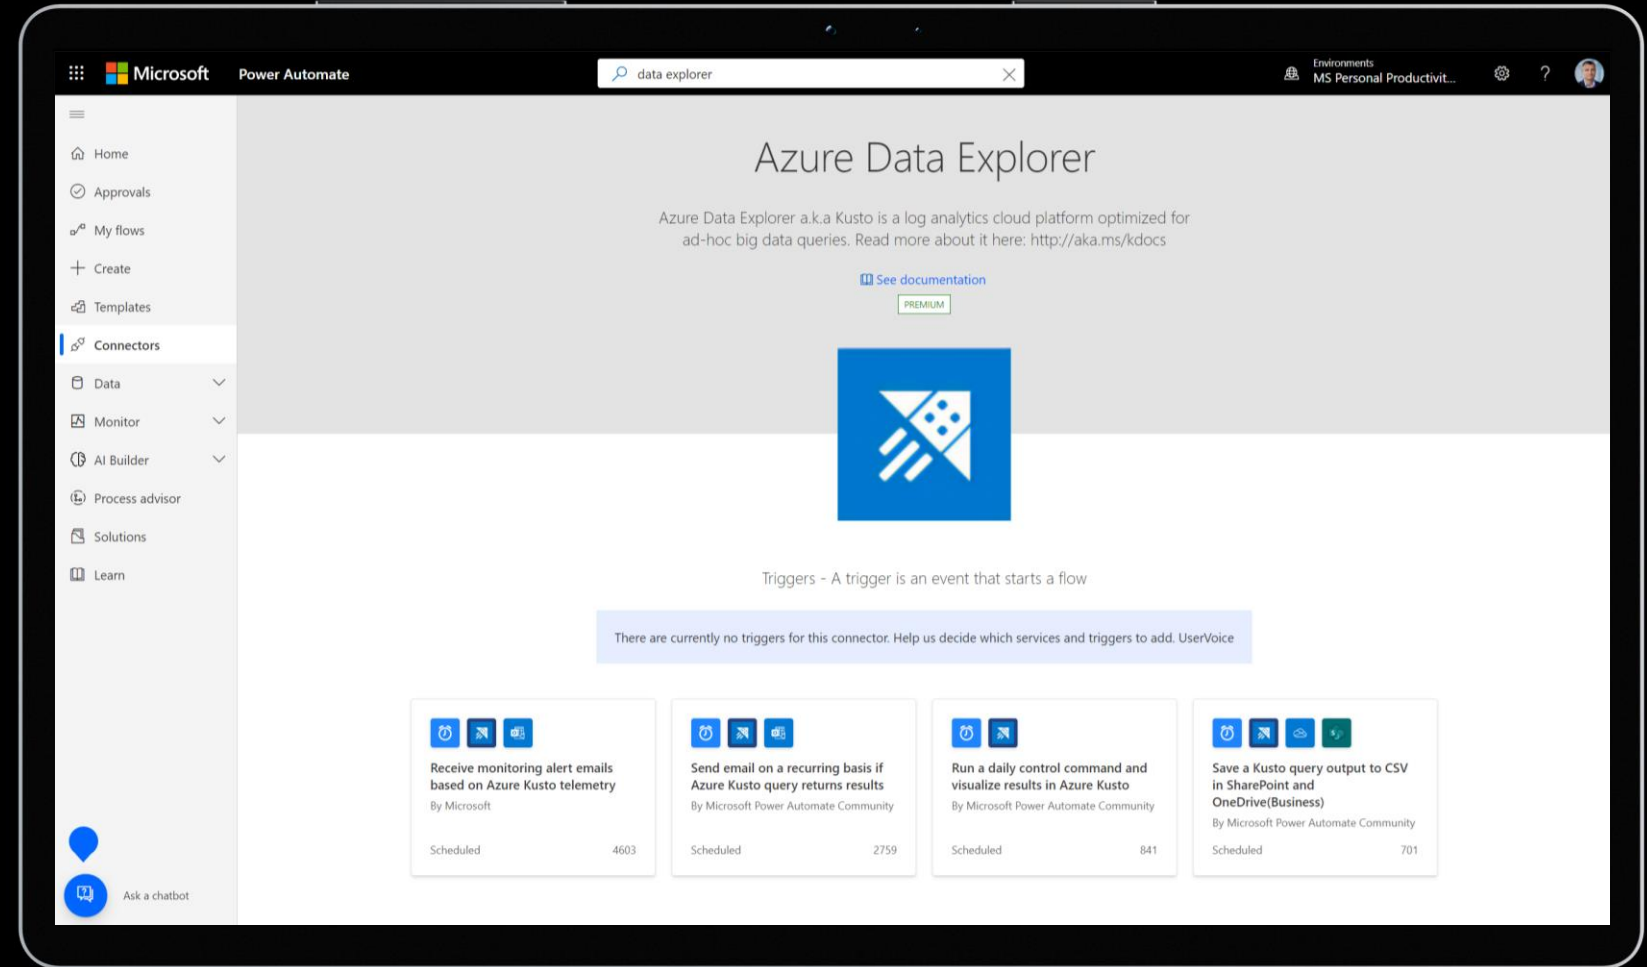
StartTime	EndTime	EpisodeId	EventId	State	EventType	InjuriesDirect	InjuriesIndirect	DeathsDirect	DeathsIndirect
> 2007-12-31 23:53:00.0000	2007-12-31 23:53:00.0000	12,037	65,839	CALIFORNIA	High Wind	0	0	0	0
> 2007-12-31 23:53:00.0000	2007-12-31 23:53:00.0000	12,037	65,838	CALIFORNIA	High Wind	0	0	0	0
> 2007-12-31 22:30:00.0000	2007-12-31 23:59:00.0000	12,950	71,588	MICHIGAN	Winter Storm	0	0	0	0
> 2007-12-31 22:30:00.0000	2007-12-31 23:59:00.0000	12,950	71,589	MICHIGAN	Winter Storm	0	0	0	0
> 2007-12-31 22:30:00.0000	2007-12-31 23:59:00.0000	12,950	71,590	MICHIGAN	Winter Storm	0	0	0	0
> 2007-12-31 20:30:00.0000	2007-12-31 23:59:00.0000	12,950	71,586	MICHIGAN	Winter Storm	0	0	0	0
> 2007-12-31 20:30:00.0000	2007-12-31 23:59:00.0000	12,950	71,587	MICHIGAN	Winter Storm	0	0	0	0
> 2007-12-31 19:00:00.0000	2007-12-31 23:59:00.0000	12,994	71,915	ALASKA	High Wind	0	0	0	0
> 2007-12-31 19:00:00.0000	2007-12-31 23:59:00.0000	12,994	71,917	ALASKA	High Wind	0	0	0	0
> 2007-12-31 18:30:00.0000	2007-12-31 23:59:00.0000	13,007	71,979	CALIFORNIA	High Wind	0	0	0	0

Connector for Power Automate, Logic Apps and Power Apps

Run KQL queries and render
charts

Send notifications, alerts,
reports

Schedule jobs using control
commands on clusters

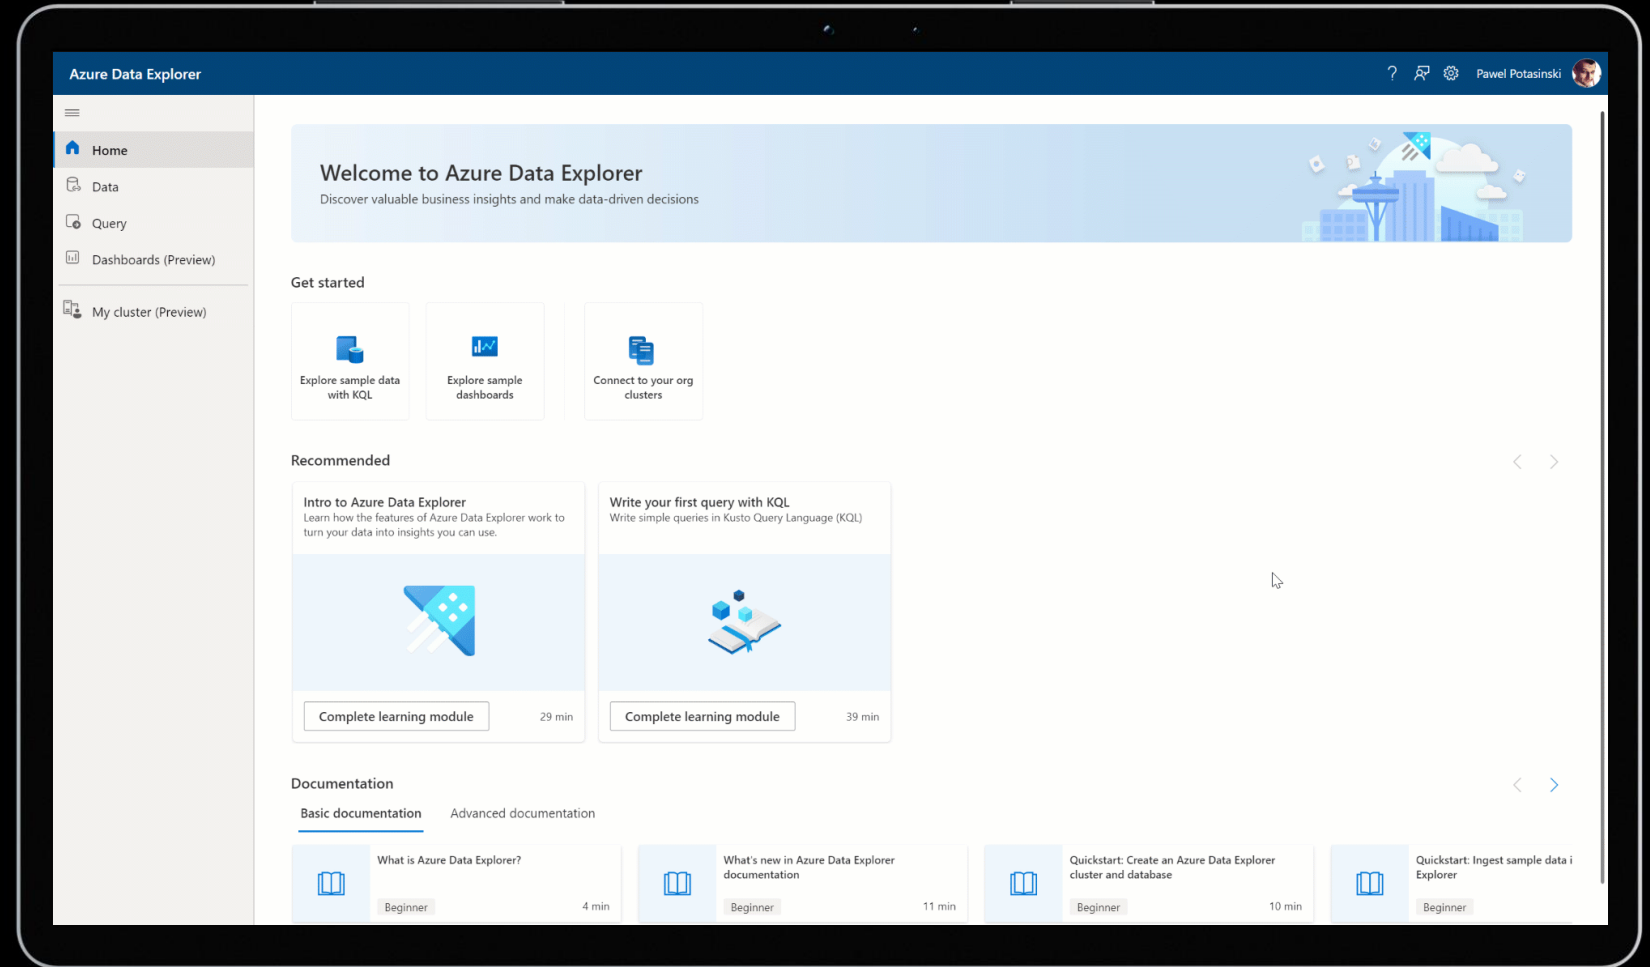


Web explorer new homepage and sample gallery

Get started with samples

Learn ADX and KQL using
recommended MS Learn
modules

Get quick access to
documentation



Start & learn for free: [Azure Data Explorer in 60 minutes](https://aka.ms/kustofree) | <https://aka.ms/kustofree>



Business Intelligence

April-June 2022

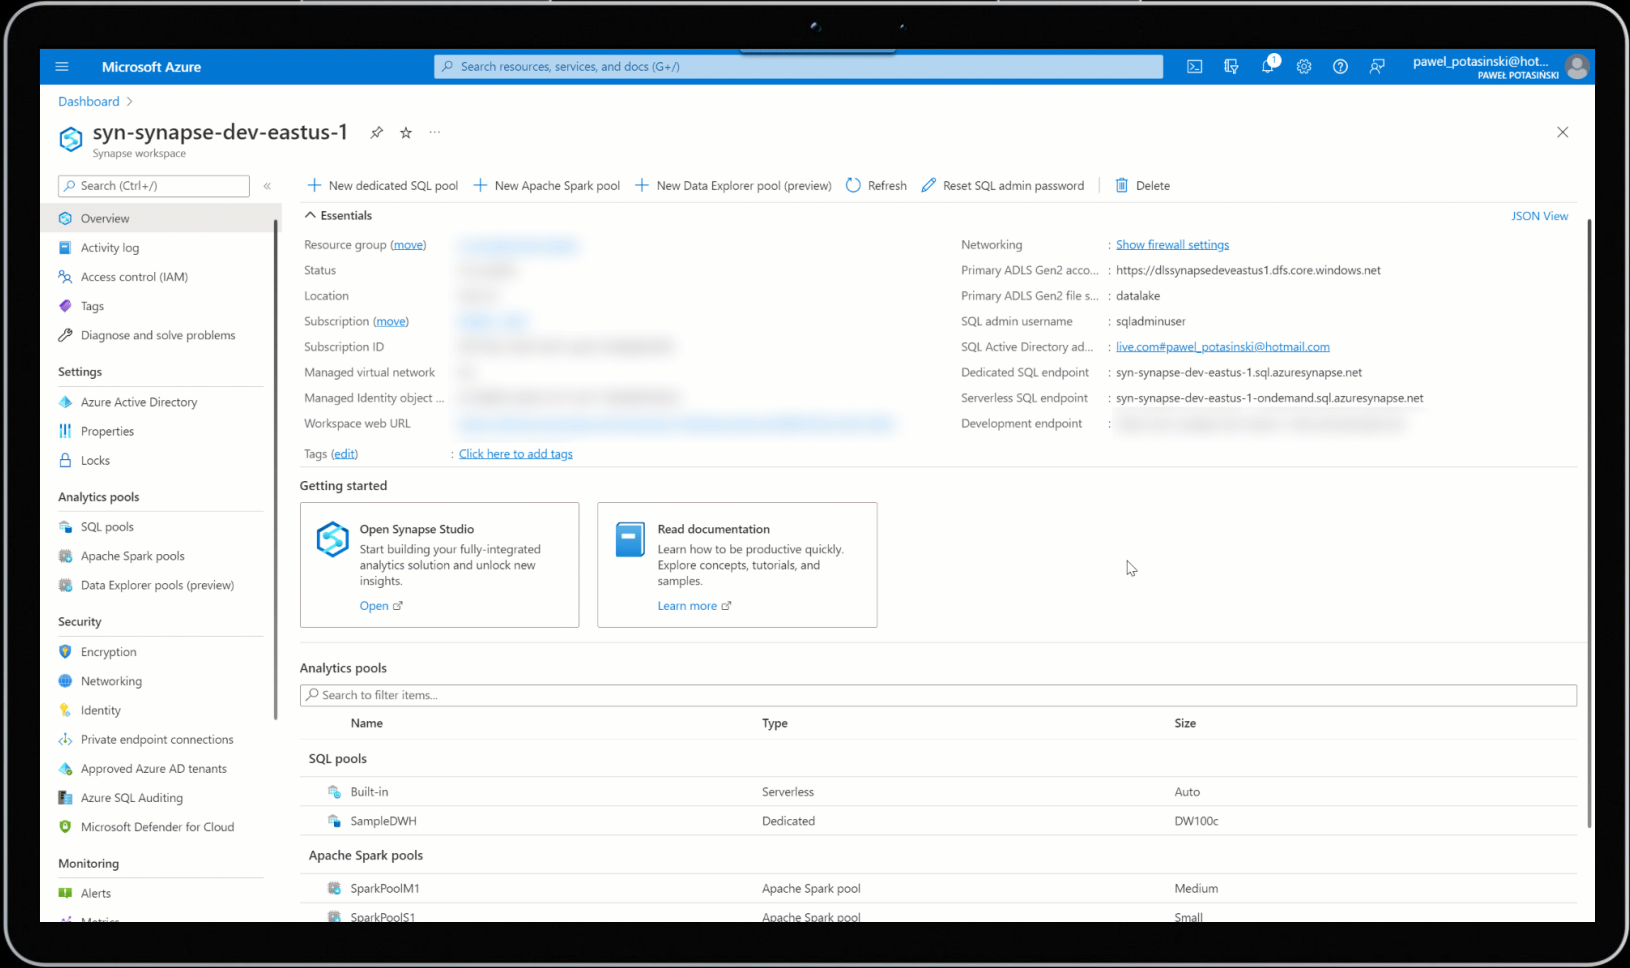
- **Public preview of datamart in Power BI**
 - General availability of automatic aggregations in Power BI
 - Power BI integration for PowerPoint (storytelling)
 - Goals in Power BI is becoming 'metrics' and is moving to GA
-
- Power BI April 2022 Feature Summary
 - Power BI May 2022 Feature Summary
 - Power BI June 2022 Feature Summary

Public Preview


Power BI
datamart

Combination of Power BI
dataflow + Azure SQL DB +
Power BI dataset

Allows model development
(relationships, RLS)



Synapse + Power BI





Azure Synapse Community Resources

Community Call to Action

- ✓ Follow us on Twitter: [@Azure_Synapse](https://twitter.com/Azure_Synapse)
- ✓ Read and comment on our blog: <https://aka.ms/SynapseBlog>
- ✓ Check monthly updates blog: <https://aka.ms/SynapseMonthlyUpdate>
- ✓ Subscribe to YouTube channel: <https://aka.ms/SynapseYouTube>
- ✓ Share ideas and vote for them: <https://aka.ms/SynapseIdeas>
- ✓ Join Azure Synapse Influencers program: <https://aka.ms/SynapseInfluencers>

Product

- Product page: <https://aka.ms/Synapse>
- Documentation: <https://aka.ms/SynapseDocs>
- Reference architectures: <https://aka.ms/SynapseArchitectures>
- Security whitepaper: <http://aka.ms/SynapseSecurity>

Technical Discussions and Q&A

- Microsoft Q&A: <https://aka.ms/SynapseQuestions>
- StackOverflow: <https://aka.ms/SynapseStackOverflow>

Quickstart

- Get started in 60 minutes: <https://aka.ms/SynapseGetStarted>
- Azure Synapse Analytics Toolkit: <https://aka.ms/SynapseToolkit>

Free Learning

- MS Learn – learning paths: <https://aka.ms/SynapseLearningPaths>
- Synapse Practitioner: <https://aka.ms/SynapsePractitioner>
- Microsoft Virtual Training Days: <https://aka.ms/SynapseMVTd>
- 30-Day Cloud Skills Challenge: <https://aka.ms/SynapseSkillsChallenge>

Samples & Accelerators

- Samples on GitHub: <https://aka.ms/SynapseSamples>
- Accelerator – End-to-End Analytics: <https://aka.ms/azsynapseee2e-git>







DATA TOBOGGAN

What's New and Cool in Azure Synapse Analytics

Pawel Potasinski
Senior Program Manager

 @pawelpotasinski
 /in/pawelpotasinski

Get this slide deck from:
<https://github.com/pawelpo/presentations>

